```
FROM continuumio/miniconda3:4.9.2

ARG port=8888

ENV NOTEBOOK_PORT $port

RUN conda config --set auto_update_conda false \
    && conda config --set notify_outdated_conda false \
    && conda config --prepend channels conda-forge \
    && conda config --set channel_priority strict \
    && conda install -Sy \
        python==3.8.5 \
        pip==20.2.4 \
        notebook=6.1.4 \
        ipywidgets=7.5.1 \
        jupyter_contrib_nbextensions=0.5.1 \
        tini=0.18.0 \
        numpy=1.19.1 \
        pandas=1.1.2 \
        matplotlib=3.2.2 \
        seaborn=0.11.0 \
    && conda clean -afy

COPY jupyter_notebook_config.py /root/.jupyter/

WORKDIR "/mnt"

ENTRYPOINT ["tini", "-g", "--"]

CMD ["jupyter", "notebook"]
```

Data analyses in Docker

1. Write a *Dockerfile*

2. Build an *image*

3. Run a *container*

Dockerfile Instructions


RUN

- Executes commands in a new "*layer*" on top of the current image and "*commits*" the result

- Docker images work a lot like git repositories

- Each instruction modifies the image by adding a new layer on top of it

- Docker stores images as a series of “diffs” between layers

- RUN creates a container from the current layer

3. Run a *container*

2. Build an *image*

1. Write a *Dockerfile*

- Executes commands in a new "*layer*" on top of the

current image and “*commits*” the result

- Docker stores images as a series of “diffs” between

new layer on top of it

layers

- Each instruction modifies the image by adding a