



# *Docker for scientific research*



Paxton Fitzpatrick  
December 1, 2021

# What is Docker?

- Self-contained, isolated software environments called “containers”
- Create, share, and deploy applications & services
- Popular tool in production environments



# But how is this useful for research?

Reproducible code!

- Experiments
- Data analyses

# But how is this useful for research?

Reproducible code!

- Experiments
  - Run the same on any computer
- Data analyses

# But how is this useful for research?

## Reproducible code!

- Experiments
  - Run the same on any computer
  - Run multiple experiments on the same computer
- Data analyses

# But how is this useful for research?

## Reproducible code!

- Experiments
  - Run the same on any computer
  - Run multiple experiments on the same computer
  - Launch complex experiments with a single command
- Data analyses

# But how is this useful for research?

## Reproducible code!

- Experiments
  - Run the same on any computer
  - Run multiple experiments on the same computer
  - Launch complex experiments with a single command
- Data analyses
  - Ensure identical packages & versions across project team

# But how is this useful for research?

## Reproducible code!

- Experiments
  - Run the same on any computer
  - Run multiple experiments on the same computer
  - Launch complex experiments with a single command
- Data analyses
  - Ensure identical packages & versions across project team
  - Isolate environments for each project from each other



# But how is this useful for research?

## Reproducible code!

- Experiments
  - Run the same on any computer
  - Run multiple experiments on the same computer
  - Launch complex experiments with a single command
- ➡ • Data analyses
  - Ensure identical packages & versions across project team
  - Isolate environments for each project from each other
  - Run analyses on Discovery in the same environment you use locally

# Quick Docker jargon...

## Image

“Template” environment with everything needed to perform a certain task

## Container

A running instance of an image

## Dockerfile

File containing instructions to build an image

## docker-compose

Tool for defining and running multiple, coordinated containers

## Docker Hub

GitHub-like website with repositories of pre-built images

# Data analyses in Docker

**Goal:** create an environment that:

- has Python installed
- contains all packages needed for analyses
- can run Jupyter notebooks
- is isolated from the host machine, but has access to it
- is easy to create, use, tweak, and share
- will always be exactly the same, no matter when or where it's used

# Data analyses in Docker

## Steps to create and use our environment:

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

# Data analyses in Docker

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

```
FROM continuumio/miniconda3:4.9.2

ARG port=8888

ENV NOTEBOOK_PORT $port

RUN conda config --set auto_update_conda false \
    && conda config --set notify_outdated_conda false \
    && conda config --prepend channels conda-forge \
    && conda config --set channel_priority strict \
    && conda install -Sy \
        python==3.8.5 \
        pip==20.2.4 \
        notebook=6.1.4 \
        ipywidgets=7.5.1 \
        jupyter_contrib_nbextensions=0.5.1 \
        tini=0.18.0 \
        numpy=1.19.1 \
        pandas=1.1.2 \
        matplotlib=3.2.2 \
        seaborn=0.11.0 \
    && conda clean -afy

COPY jupyter_notebook_config.py /root/.jupyter/

WORKDIR "/mnt"

ENTRYPOINT ["tini", "-g", "--"]

CMD ["jupyter", "notebook"]
```

# Data analyses in Docker

## Dockerfile Instructions

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

### FROM

- Specifies the base image to build on top of
- Can be an image you built or any image on Docker Hub
- Format:  
`FROM <user>/<repository>:<tag>`
- Debian 10 ("Buster") image with Miniconda 4.9.2 installed

```
FROM continuumio/miniconda3:4.9.2
```

```
ARG port=8888
```

```
ENV NOTEBOOK_PORT $port
```

# Data analyses in Docker

## Dockerfile Instructions

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

```
FROM debian:buster-slim

LABEL maintainer="Anaconda, Inc"

ENV LANG=C.UTF-8 LC_ALL=C.UTF-8

RUN apt-get update -q && \
    apt-get install -q -y --no-install-recommends \
        bzip2 \
        ca-certificates \
        git \
        libglib2.0-0 \
        libsm6 \
        libxext6 \
        libxrender1 \
        mercurial \
        openssh-client \
        subversion \
        wget \
    && apt-get clean \
    && rm -rf /var/lib/apt/lists/*

ENV PATH /opt/conda/bin:$PATH

CMD [ "/bin/bash" ]

ARG CONDA_VERSION=py38_4.9.2

RUN set -x && \
    UNAME_M="$(uname -m)" && \
    if [ "${UNAME_M}" = "x86_64" ]; then \
        MINICONDA_URL="https://repo.anaconda.com/miniconda/Miniconda3-${CONDA_VERSION}-Linux-x86_64.sh"; \
        SHA256SUM="1314b90489f154602fd794accfc90446111514a5a72fe1f71ab83e07de9504a7"; \
    elif [ "${UNAME_M}" = "s390x" ]; then \
        MINICONDA_URL="https://repo.anaconda.com/miniconda/Miniconda3-${CONDA_VERSION}-Linux-s390x.sh"; \
        SHA256SUM="4e6ace66b732170689fd2a7d86559f674f2de0a0a0fbaefd86ef597d52b89d16"; \
    elif [ "${UNAME_M}" = "aarch64" ]; then \
        MINICONDA_URL="https://repo.anaconda.com/miniconda/Miniconda3-${CONDA_VERSION}-Linux-aarch64.sh"; \
        SHA256SUM="b6fbb97d7cef35ebee8739536752cd8b8b414f88e237146b11ebf081c44618f"; \
    elif [ "${UNAME_M}" = "ppc64le" ]; then \
        MINICONDA_URL="https://repo.anaconda.com/miniconda/Miniconda3-${CONDA_VERSION}-Linux-ppc64le.sh"; \
        SHA256SUM="2b111dab4b72a34c969188aa7a91eca927a034b14a87f725fa8d295955364e71"; \
    fi && \
    wget "${MINICONDA_URL}" -O miniconda.sh -q && \
    echo "${SHA256SUM} miniconda.sh" > shasum && \
    if [ "${CONDA_VERSION}" != "latest" ]; then sha256sum --check --status shasum; fi && \
    mkdir -p /opt && \
    sh miniconda.sh -b -p /opt/conda && \
    rm miniconda.sh shasum && \
    ln -s /opt/conda/etc/profile.d/conda.sh /etc/profile.d/conda.sh && \
    echo ". /opt/conda/etc/profile.d/conda.sh" >> ~/.bashrc && \
    echo "conda activate base" >> ~/.bashrc && \
    find /opt/conda/ -follow -type f -name '*.a' -delete && \
    find /opt/conda/ -follow -type f -name '*.js.map' -delete && \
    /opt/conda/bin/conda clean -afy
```

```
FROM continuumio/miniconda3:4.9.2
```

```
ARG port=8888
```

```
ENV NOTEBOOK_PORT $port
```

# Data analyses in Docker

## Dockerfile Instructions

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

### ARG

- Defines an argument that can be passed via the command line when building the image
- The port the Jupyter notebook server will listen on
- Variable persists for the remainder of the build

```
FROM continuumio/miniconda3:4.9.2
```

```
ARG port=8888
```

```
ENV NOTEBOOK_PORT $port
```

```
RUN conda config --set auto_update_conda false \  
&& conda config --set notify_outdated_conda false \  
&& conda config --set auto_update_conda false \  
&& conda config --set notify_outdated_conda false \  
&& conda config --set auto_update_conda false \  
&& conda config --set notify_outdated_conda false
```



# Data analyses in Docker

## Dockerfile Instructions

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

### ARG

- Defines an argument that can be passed via the command line when building the image
- The port the Jupyter notebook server will listen on
- Variable persists for the remainder of the build

### ENV

- Sets an environment variable for the remainder of the build *and* in containers run from the resulting image

```
FROM continuumio/miniconda3:4.9.2
```

```
ARG port=8888
```

```
ENV NOTEBOOK_PORT $port
```

```
RUN conda config --set auto_update_conda false \
    && conda config --set notify_outdated_conda false \
    && conda config --prepend channels conda-forge \
    && conda config --set channel_priority strict \
```

# Data analyses in Docker

## Dockerfile Instructions

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

### RUN

- Executes commands in a new “*layer*” on top of the current image and “*commits*” the result
- **Docker images work a lot like git repositories**
  - Each instruction modifies the image by adding a new layer on top of it
  - Docker stores images as a series of “*diffs*” between layers
- RUN creates a container from the current layer

```
ARG port=8888
```

```
ENV NOTEBOOK_PORT $port
```

```
RUN conda config --set auto_update_conda false \  
&& conda config --set notify_outdated_conda false \  
&& conda config --prepend channels conda-forge \  
&& conda config --set channel_priority strict \  
&& conda install -Sy \  
    python==3.8.5 \  
    pip==20.2.4 \  
    notebook=6.1.4 \  
    ipywidgets=7.5.1 \  
    jupyter_contrib_nbextensions=0.5.1 \  
    tini=0.18.0 \  
    numpy=1.19.1 \  
    pandas=1.1.2 \  
    matplotlib=3.2.2 \  
    seaborn=0.11.0 \  
&& conda clean -afy
```

```
COPY jupyter_notebook_config.py /root/.jupyter/
```

```
WORKDIR "/mnt"
```

# Data analyses in Docker

## Dockerfile Instructions

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

### RUN

- Configure conda
  - Disable auto-updates and update notification
  - Use conda-forge channel to install packages
- Install requirements
  - Python (check expected version)
  - pip package manager
  - Jupyter notebooks, interactive widgets, some handy notebook extensions
  - `tini` init manager
  - Packages for analyses
- Clear caches to reduce image size

```
ARG port=8888
```

```
ENV NOTEBOOK_PORT $port
```

```
RUN conda config --set auto_update_conda false \
    && conda config --set notify_outdated_conda false \
    && conda config --prepend channels conda-forge \
    && conda config --set channel_priority strict \
    && conda install -Sy \
        python==3.8.5 \
        pip==20.2.4 \
        notebook=6.1.4 \
        ipywidgets=7.5.1 \
        jupyter_contrib_nbextensions=0.5.1 \
        tini=0.18.0 \
        numpy=1.19.1 \
        pandas=1.1.2 \
        matplotlib=3.2.2 \
        seaborn=0.11.0 \
    && conda clean -afy
```

```
COPY jupyter_notebook_config.py /root/.jupyter/
```

```
WORKDIR "/mnt"
```

# Data analyses in Docker

## Dockerfile Instructions

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

### RUN

- Why all the **&&**'s?

```
ARG port=8888
```

```
ENV NOTEBOOK_PORT $port
```

```
RUN conda config --set auto_update_conda false \  
&& conda config --set notify_outdated_conda false \  
&& conda config --prepend channels conda-forge \  
&& conda config --set channel_priority strict \  
&& conda install -Sy \  
    python==3.8.5 \  
    pip==20.2.4 \  
    notebook=6.1.4 \  
    ipywidgets=7.5.1 \  
    jupyter_contrib_nbextensions=0.5.1 \  
    tini=0.18.0 \  
    numpy=1.19.1 \  
    pandas=1.1.2 \  
    matplotlib=3.2.2 \  
    seaborn=0.11.0 \  
&& conda clean -afy
```

```
COPY jupyter_notebook_config.py /root/.jupyter/
```

```
WORKDIR "/mnt"
```



# Data analyses in Docker

## Dockerfile Instructions

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

### RUN

- Why all the **&&**'s?
- Each RUN instruction creates an intermediate layer

2.19GB

```
RUN conda config --set auto_update_conda false
RUN conda config --set notify_outdated_conda false
RUN conda config --prepend channels conda-forge
RUN conda config --set channel_priority strict
RUN conda install -Sy \
    python==3.8.5 \
    pip==20.2.4
RUN conda install -Sy \
    notebook=6.1.4 \
    ipywidgets=7.5.1 \
    jupyter_contrib_nbextensions=0.5.1 \
    tini=0.18.0
RUN conda install -Sy \
    numpy=1.19.1 \
    pandas=1.1.2 \
    matplotlib=3.2.2 \
    seaborn=0.11.0
RUN conda clean -afy
```

ARG port=8888

ENV NOTEBOOK\_PORT \$port 918MB

```
RUN conda config --set auto_update_conda false \
    && conda config --set notify_outdated_conda false \
    && conda config --prepend channels conda-forge \
    && conda config --set channel_priority strict \
    && conda install -Sy \
        python==3.8.5 \
        pip==20.2.4 \
        notebook=6.1.4 \
        ipywidgets=7.5.1 \
        jupyter_contrib_nbextensions=0.5.1 \
        tini=0.18.0 \
        numpy=1.19.1 \
        pandas=1.1.2 \
        matplotlib=3.2.2 \
        seaborn=0.11.0 \
    && conda clean -afy
```

COPY jupyter\_notebook\_config.py /root/.jupyter/

# Data analyses in Docker

## Dockerfile Instructions

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

### RUN

- Why all the **&&**'s?
- Each RUN instruction creates an intermediate layer

```
ARG port=8888
```

```
ENV NOTEBOOK_PORT $port
```

```
RUN conda config --set auto_update_conda false \  
&& conda config --set notify_outdated_conda false \  
&& conda config --prepend channels conda-forge \  
&& conda config --set channel_priority strict \  
&& conda install -Sy \  
    python==3.8.5 \  
    pip==20.2.4 \  
    notebook=6.1.4 \  
    ipywidgets=7.5.1 \  
    jupyter_contrib_nbextensions=0.5.1 \  
    tini=0.18.0 \  
    numpy=1.19.1 \  
    pandas=1.1.2 \  
    matplotlib=3.2.2 \  
    seaborn=0.11.0 \  
&& conda clean -afy
```

```
COPY jupyter_notebook_config.py /root/.jupyter/
```

```
WORKDIR "/mnt"
```

# Data analyses in Docker

## Dockerfile Instructions

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

### COPY

- Copies files or directories from the host into the image
- Config for Jupyter notebook server

```
&& conda config --prepend channels conda-forge \  
&& conda config --set channel_priority strict \  
&& conda install -Sy \  
    python==3.8.5 \  
    pip==20.2.4 \  
    notebook=6.1.4 \  
    ipywidgets=7.5.1 \  
    jupyter_contrib_nbextensions=0.5.1 \  
    tini=0.18.0 \  
    numpy=1.19.1 \  
    pandas=1.1.2 \  
    matplotlib=3.2.2 \  
    seaborn=0.11.0 \  
&& conda clean -afy
```

```
COPY jupyter_notebook_config.py /root/.jupyter/
```

```
WORKDIR "/mnt"
```

```
ENTRYPOINT ["tini", "-g", "--"]
```

# Data analyses in Docker

## Dockerfile Instructions

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

### WORKDIR

- Sets the working directory for subsequent instructions and containers run from the image
- /mnt will be the “mount point” for our container

```
&& conda install -Sy \
    python==3.8.5 \
    pip==20.2.4 \
    notebook=6.1.4 \
    ipywidgets=7.5.1 \
    jupyter_contrib_nbextensions=0.5.1 \
    tini=0.18.0 \
    numpy=1.19.1 \
    pandas=1.1.2 \
    matplotlib=3.2.2 \
    seaborn=0.11.0 \
&& conda clean -afy
```

```
COPY jupyter_notebook_config.py /root/.jupyter/
```

```
WORKDIR "/mnt"
```

```
ENTRYPOINT ["tini", "-g", "--"]
```

```
CMD ["jupyter", "notebook"]
```



# Data analyses in Docker

## Dockerfile Instructions

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

### ENTRYPOINT

- Executable *always* run when a container starts.  
*Cannot* be overridden from the command line.

### CMD

- *Default* arguments passed to ENTRYPOINT (or default command to run, if no ENTRYPOINT). *Can* be overridden from the command line.

```
python=3.7.4 \
notebook=6.1.4 \
ipywidgets=7.5.1 \
jupyter_contrib_nbextensions=0.5.1 \
tini=0.18.0 \
numpy=1.19.1 \
pandas=1.1.2 \
matplotlib=3.2.2 \
seaborn=0.11.0 \
&& conda clean -afy
```

```
COPY jupyter_notebook_config.py /root/.jupyter/
```

```
WORKDIR "/mnt"
```

```
ENTRYPOINT ["tini", "-g", "--"]
```

```
CMD ["jupyter", "notebook"]
```

# Data analyses in Docker

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

```
FROM continuumio/miniconda3:4.9.2

ARG port=8888

ENV NOTEBOOK_PORT $port

RUN conda config --set auto_update_conda false \
    && conda config --set notify_outdated_conda false \
    && conda config --prepend channels conda-forge \
    && conda config --set channel_priority strict \
    && conda install -Sy \
        python==3.8.5 \
        pip==20.2.4 \
        notebook=6.1.4 \
        ipywidgets=7.5.1 \
        jupyter_contrib_nbextensions=0.5.1 \
        tini=0.18.0 \
        numpy=1.19.1 \
        pandas=1.1.2 \
        matplotlib=3.2.2 \
        seaborn=0.11.0 \
    && conda clean -afy

COPY jupyter_notebook_config.py /root/.jupyter/

WORKDIR "/mnt"

ENTRYPOINT ["tini", "-g", "--"]

CMD ["jupyter", "notebook"]
```

# Data analyses in Docker

## The `docker build` command

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

- Builds an image from a Dockerfile and a "*build context*"
- "Build context": files that can be used during the build process
- Can be a local directory path, GitHub repo URL, or tarball (`.tar.gz` file)
- Default: pass context to `docker build`; Docker looks for "Dockerfile" in its root directory

# Data analyses in Docker

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

## The `docker build` command

- To build the image from our Dockerfile:

```
$ docker build -t tutorial-image .
```

Name and optionally tag the  
image (default tag: "latest")

Build context: current directory

(If you're following along, don't run this)

# Data analyses in Docker

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

## The `docker build` command

- To build the image from our Dockerfile:

```
$ docker build -t tutorial-image --build-arg port=8889 .
```

Set the `$port` variable from the ARG instruction in the Dockerfile

Name and optionally tag the image (default tag: "latest")

Build context: current directory

(If you're following along, don't run this)

# Data analyses in Docker

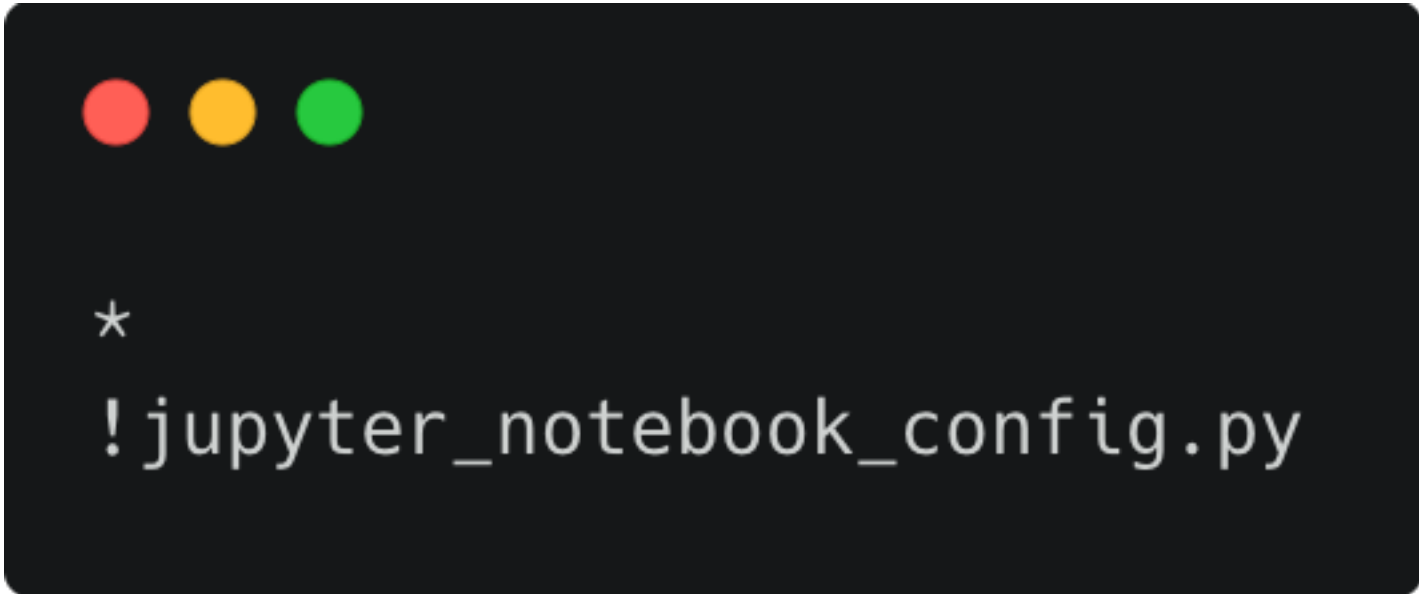
1. Write a *Dockerfile*

2. Build an *image*

3. Run a *container*

## The `docker build` command

- Large build context can lead to very slow build
- You can exclude files and directories from context with a `.dockerignore` file



```
*  
!jupyter_notebook_config.py
```

A terminal window with a dark background and three colored window control buttons (red, yellow, green) in the top-left corner. It displays two lines of text: an asterisk on the first line and '!jupyter\_notebook\_config.py' on the second line.

# Data analyses in Docker

1. Write a *Dockerfile*
2. Build an *image*
3. Run a *container*

## The `docker run` command

- `docker run` creates and runs a container from an image
- Can be an image you've built locally or any image on Docker Hub
- To run a container from the pre-built version of our image on Docker Hub:

Make the container interactive  
via the terminal

Assign a name to the container

Specify the image used to run the container

```
$ docker run -it -p 8888:8888 --name analyses -v $PWD:/mnt paxtonfitzpatrick/tutorial-image
```

Mount the current working directory  
into `/mnt` in the container

Bind port 8888 in the container to port 8888 on the  
host. Allows us to access notebook server in  
container from browser on host

# Data analyses in Docker

Demo... 🙌




# Data analyses in Docker

Using our container on Discovery

Demo... 🙌

# Behavioral experiments in Docker

Running psiTurk through Docker



```
FROM python:3.6-stretch

ENV PSITURK_GLOBAL_CONFIG_LOCATION "/exp"

WORKDIR "/exp"



RUN pip install \
    psiturk==2.3.8 \
    pymysql==0.10.0 \
    python-Levenshtein==0.12.0 \
    && rm -rf ~/.cache/pip

CMD ["bash"]
```

# Behavioral experiments in Docker

## Running psiTurk through Docker

- Specify the base image to build from
- Set environment variable (tells psiTurk where to look for `.psiturkconfig`)
- Set the working director
- Install psiTurk and some extra required packages
- Clear cache to keep image small
- Set the command executed when running a container (launch a bash shell)



```
FROM python:3.6-stretch

ENV PSITURK_GLOBAL_CONFIG_LOCATION "/exp"

WORKDIR "/exp"

RUN pip install \
    psiturk==2.3.8 \
    pymysql==0.10.0 \
    python-Levenshtein==0.12.0 \
    && rm -rf ~/.cache/pip

CMD ["bash"]
```

# Behavioral experiments in Docker

## Running psiTurk through Docker

- We need 3 support services to run on MTurk
- Normally complex to configure, tedious to manage
- But Docker makes this *much* easier!
  - All 3 available on Docker Hub as ready-to-run images
  - docker-compose: build, start and stop entire setup with a single command
- Configure full application in docker-compose.yml

```
version: '3'
services:

  psiturk:
    container_name: my-experiment
    build: .
    volumes:
      - ./exp:/exp
    tty: true
    stdin_open: true
    restart: unless-stopped

  nginx:
    container_name: my-experiment-nginx
    image: nginx:latest
    ports:
      - 80:80
    volumes:
      - ./exp:/var/www/exp:ro
      - ./default.conf:/etc/nginx/conf.d/default.conf
    restart: unless-stopped

  db:
    container_name: my-experiment-db
    image: mysql:5.7
    volumes:
      - ./data/db:/var/lib/mysql
    environment:
      MYSQL_ROOT_PASSWORD: mypassword
      MYSQL_DATABASE: participants
      MYSQL_USER: paxton
      MYSQL_PASSWORD: psiturk
    restart: unless-stopped


  adminer:
    container_name: my-experiment-adminer
    image: adminer:latest
    ports:
      - 127.0.0.1:8080:8080
```

# Behavioral experiments in Docker

`docker-compose.yml`

## Container 1: **psiTurk server**

- Build context
- Mount the host directory `exp` to `/exp` in the container
- Allocate a pseudo-TTY and send stdin to the container (so we can run an interactive shell)
- If the container fails for some reason, restart it unless we explicitly stopped it.



```
version: '3'
services:

  psiturk:
    container_name: my-experiment
    build: .
    volumes:
      - ./exp:/exp
    tty: true
    stdin_open: true
    restart: unless-stopped

  nginx:
    container_name: my-experiment-nginx
    image: nginx:latest
    ports:
```

# Behavioral experiments in Docker

`docker-compose.yml`

## Container 2: NGINX

(reverse proxy server for load balancing)

- Pull the `nginx:latest` image from Docker Hub
- Map host port 80 to container port 80
- Create two mount points:
  - host `exp` directory ➡ NGINX web root
  - NGINX config file ➡ expected location



```
build: .  
volumes:  
  - ./exp:/exp  
tty: true  
stdin_open: true  
restart: unless-stopped
```

```
nginx:  
  container_name: my-experiment-nginx  
  image: nginx:latest  
  ports:  
    - 80:80  
  volumes:  
    - ./exp:/var/www/exp:ro  
    - ./default.conf:/etc/nginx/conf.d/default.conf  
  restart: unless-stopped
```

```
db:  
  container_name: my-experiment-db  
  image: mysql:5.7  
  volumes:  
    - ./data/db:/var/lib/mysql  
  environment:
```




# Behavioral experiments in Docker

`docker-compose.yml`

## Container 3: MySQL database

(greater concurrency than default SQLite DB)

- Pull the `mysql:5.7` image from Docker Hub
- Mount `data/db` directory on host to `/var/lib/mysql` in container
- Set container environment variables:
  - Password for MySQL root user (required)
  - Name of DB created in the container
  - Username/password for psiTurk to write to DB



```
- 80:80
volumes:
- ./exp:/var/www/exp:ro
- ./default.conf:/etc/nginx/conf.d/default.conf
restart: unless-stopped
```

```
db:
  container_name: my-experiment-db
  image: mysql:5.7
  volumes:
  - ./data/db:/var/lib/mysql
  environment:
    MYSQL_ROOT_PASSWORD: mypassword
    MYSQL_DATABASE: participants
    MYSQL_USER: paxton
    MYSQL_PASSWORD: psiturk
  restart: unless-stopped
```

```
adminer:
  container_name: my-experiment-adminer
  image: adminer:latest
  ports:
  - 127.0.0.1:8080:8080
```

# Behavioral experiments in Docker

`docker-compose.yml`

## Container 4: Adminer

(PHP app for viewing/downloading data)

- Pull the `adminer:latest` image from Docker Hub
- Map host port 8080 to container port 8080 (allows us to access Adminer via a web browser)

```
volumes:  
  - ./data/db:/var/lib/mysql  
environment:  
  MYSQL_ROOT_PASSWORD: mypassword  
  MYSQL_DATABASE: participants  
  MYSQL_USER: paxton  
  MYSQL_PASSWORD: psiturk  
restart: unless-stopped
```

```
adminer:  
  container_name: my-experiment-adminer  
  image: adminer:latest  
  ports:  
    - 127.0.0.1:8080:8080
```



# Behavioral experiments in Docker

```
version: '3'
services:

  psiturk:
    container_name: my-experiment
    build: .
    volumes:
      - ./exp:/exp
    tty: true
    stdin_open: true
    restart: unless-stopped

  nginx:
    container_name: my-experiment-nginx
    image: nginx:latest
    ports:
      - 80:80
    volumes:
      - ./exp:/var/www/exp:ro
      - ./default.conf:/etc/nginx/conf.d/default.conf
    restart: unless-stopped

  db:
    container_name: my-experiment-db
    image: mysql:5.7
    volumes:
      - ./data/db:/var/lib/mysql
    environment:
      MYSQL_ROOT_PASSWORD: mypassword
      MYSQL_DATABASE: participants
      MYSQL_USER: paxton
      MYSQL_PASSWORD: psiturk
    restart: unless-stopped

  adminer:
    container_name: my-experiment-adminer
    image: adminer:latest
    ports:
      - 127.0.0.1:8080:8080
```

# Behavioral experiments in Docker

Demo... 🙌